



COP 3540: Introduction to Database Structures

Fall 2017

Query Optimization III

Query Processing

- **access paths** are alternative ways to retrieve tuples from a relation
- two kinds of access paths:
 1. file scan
 2. index plus a matching selection condition
- **selectivity** of an access path is the number of pages retrieved (index pages + data pages) if index access path is used to retrieve all desired tuples
- **most selective** access path is the one that retrieves the fewest pages
- using the most selective access path minimizes the cost of data retrieval

Query Processing

The following schema will be used:

Sailors(sid: integer, sname: string, rating: integer, age: real)

Reserves(sid: integer, bid: integer, day: dates, rname: string)

rname has been added to reserves (reservation may be made by a person who is not a sailor)

Assumptions:

Reserves:

- a tuple is 40 bytes long
- a page can hold 100 tuples
- 1,000 pages

Sailors:

- a tuple is 50 bytes long
- a page can hold 80 tuples
- 500 pages

Query Processing

We will:

- consider only I/O costs
- measure I/O cost in terms of the # of page I/Os
- use big-O notation to express the complexity of an algorithm in terms of an input parameter (cost of a file scan is $O(M)$, where M is the size of the file)

Query Processing

Selection Operation:

```
SELECT *  
FROM   Reserves R  
WHERE  R.rname='Joe'
```

$\sigma R.attr \text{ op } value(R)$

1. No Index, Unsorted Data

- scan the entire relation
- check the condition on each tuple ($R.attr \text{ op } value$)
- add the tuple to the result if the condition is satisfied
- cost is M I/Os (M is # of pages) = 1,000 I/Os (Reserves contains 1,000 pages)

expensive because it does not utilize the selection to reduce the number of tuples retrieved in any way!

Query Processing

Selection Operation:

2. No Index, Sorted Data (R is physically sorted on $R.attr$)

- sorted-file scan with selection condition $\sigma R.attr \text{ op value}(R)$
- binary search to locate the first tuple that satisfies the selection condition ($R.attr1 > 5$, and that R is sorted on $attr1$ in ascending order)
- retrieve all tuples that satisfy the selection condition starting at this location
- cost of binary search is $O(\log_2 M)$
- cost of the scan to retrieve qualifying tuples vary from zero to M
- cost of binary search is $\log_2 1,000 \approx 10$ I/Os

Query Processing

Selection Operation:

3. B+ Tree Index

- cost of identifying the starting leaf page for the scan is typically 2 to 3 I/Os
- cost of scanning the leaf level page for qualifying data entries depends on the # of entries
- cost of retrieving qualifying tuples from R depends on:
 1. # of qualifying tuples
 2. if index is clustered
 - clustered: cost is probably just one page I/O (all tuples are contained in a single page). estimating that roughly 10 percent of Reserves tuples are in the result, a clustered B+ tree index on the rname field of Reserves, would retrieve the qualifying tuples with 100 I/Os
 - unclustered: each index entry could point to a qualifying tuple on a different page resulting in 10,000 tuples, or 100 pages and cost would be 10,000 I/Os

Query Processing

Selection Operation:


4. Hash Index, Equality Selection

- cost includes 1 or 2 I/Os to retrieve the bucket page in the index + cost of retrieving qualifying tuples from R
- unclustered hash index on the *rname* attribute:
 - 10 buffer pages
 - 100 reservations made by people named Joe
- cost of retrieving the index page containing the rids is 1 or 2 I/Os
- cost of retrieving the 100 Reserves tuples varies between 1 and 100
- If these 100 records are contained in 5 pages of Reserves = 5 I/Os (if rids are sorted by page component)

Query Processing

Projection Operation:

```
SELECT DISTINCT R.sid, R.bid  
FROM   Reserves R
```

 $\pi_{sid,bid}Reserves$

To implement projection:

1. remove unwanted attributes (those not specified in the projection)
2. eliminate any duplicate tuples that are produced by
 - sorting algorithm or
 - hashing algorithm

Query Processing

Projection Operation:

1. Projection Based on Sorting

sorting algorithm has the following steps:

1. scan R and produce a set of tuples that contain only the desired attributes, cost = M I/Os (scan R) + T I/Os (write the temporary relation T is $O(M)$) so scan is 1,000 I/Os, assume T tuple is 10 bytes, cost of writing is 250 I/Os
2. sort this set of tuples using the combination of all its attributes as the key for sorting, cost = $O(T \log T)$ (also $O(M \log M)$), if we have 20 buffer pages, we sort in two passes at a cost of $2 * 2 * 250 = 1,000$ I/Os
3. scan the sorted result, comparing adjacent tuples, and discard duplicates, cost = T (250 I/Os)

total cost is $O(M \log M) = 2,500$ I/Os

Query Processing

Projection Operation:

1. Projection Based on Sorting

improvements:

- project out unwanted attributes during the first pass (Pass 0) of sorting
- eliminate duplicates during the merging passes (fewer tuples are written out in each pass where most of the duplicates will be eliminated in the very first merging pass)

first pass: scan R cost = 1,000 I/Os + write out 250 pages
with 20 buffer pages, the 250 pages are written out as 7
internally sorted runs each about 40 pages long

second pass: read the runs cost = 250 I/Os, and merge them

total cost = 1,500 I/Os

Query Processing

Projection Operation:

2. Projection Based on Hashing

good for large number of buffer pages (B) relative to the number of pages of R

has the following phases:

1. partitioning
2. duplicate elimination

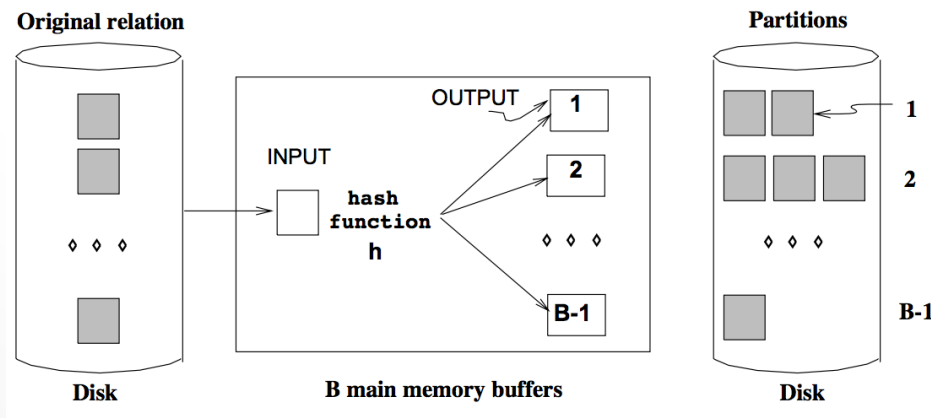
Query Processing

Projection Operation:

2. Projection Based on Hashing

1. partitioning

- one input buffer page and $B - 1$ output buffer pages
- relation R is read into the input buffer page, one page at a time
- project out unwanted attributes for each tuple
- apply a hash function h to the combination of all remaining attributes
- tuple is written to the output buffer page that it is hashed to by h
- $B - 1$ partitions at the end
- tuples that belong to different partitions are not duplicates because they have different hash values (duplicates are in the same partition)



Query Processing

Projection Operation:

2. Projection Based on Hashing

2. duplicate elimination

for each partition

1.
 - read in the partition one page at a time
 - hash each tuple by applying hash function h_2 ($\neq h_1$) to the combination of all fields
 - insert it into in-memory hash table
 - if new tuple hashes to the same value as existing tuple, compare the two to check whether the new tuple is a duplicate
 - discard duplicates as detected
 2.
 - write the duplicate-free tuples in the hash table to the result file
 - clear the in-memory hash table to prepare for the next partition
- cost: read R and write T cost = $M + T$ I/Os, cost of hashing is a CPU cost (not considered), cost of phase 2 part 2 is $M + 2T$ I/Os
 - Total cost = $1,000 + 2 * 250 = 1,500$ I/Os

Query Processing

Join Operation

SELECT *
FROM Reserves R, Sailors S
WHERE R.sid = S.sid

—————→ $R \bowtie S$

- join can be defined as a cross-product followed by selections and projections
- result of a cross-product is typically much larger than the result of a join
- recognize joins and implement them without materializing the underlying cross-product
- using two relations R and S, with the join condition $R_i = S_j$ we assume:
 - M pages in R with p_R tuples per page
 - N pages in S with p_S tuples per page

and we look at:

1. Sort-Merge Join
2. Block Nested Loops Join
3. Index Nested Loops Join
4. Sort-Merge Join
5. Hash Join

Query Processing

Join Operation

1. Sort-Merge Join

SELECT *	foreach tuple $r \in R$ do
FROM Reserves R, Sailors S	foreach tuple $s \in S$ do
WHERE R.sid = S.sid	if $r_i == s_j$ then add $\langle r, s \rangle$ to result

- cost of scanning R = M I/Os
- scan S a total of $pR * M$ times
- each scan costs N I/Os.
- total cost = $M + pR * M * N$
- $M = 1,000$, $pR = 100$, and N is 500
- total cost of simple nested loops join = $1,000 + 100 * 1,000 * 500$ page I/Os = $1,000 + (5 * 10^7)$ I/Os (huge!)

Improvement - join page-at-a-time with an improvement of a factor of pR (total cost $M + M * N = 1,000 + 1,000 * 500 = 501,000$ I/Os)

Query Processing

Join Operation

2. Block Nested Loops Join

```
SELECT *  
FROM   Reserves R, Sailors S  
WHERE  R.sid = S.sid
```

if we have enough memory to hold the smaller relation, say R, with at least two extra buffer pages left over

- each tuple $s \in S$, we check R and output a tuple $\langle r, s \rangle$ for qualifying tuples s ($r_i = s_j$), extra buffer is an output buffer
- each relation is scanned once, total I/O cost of $M + N$

Query Processing

Join Operation

2. Block Nested Loops Join

if not enough memory:

- break R into blocks that can fit into the buffer pages and scan all of S for each block of R
- R is the outer relation (scanned once), S is the inner relation (scanned multiple times)
- B buffer pages, we can read in $B - 2$ pages of the outer relation R and scan the inner relation S using one of the two remaining pages
- write out tuples $\langle r, s \rangle$, where $r \in R\text{-block}$ and $s \in S\text{-page}$ and $r_i = s_j$, using the last buffer page for output
- a way to find matching pairs of tuples is to build a main-memory hash table for the block of R .

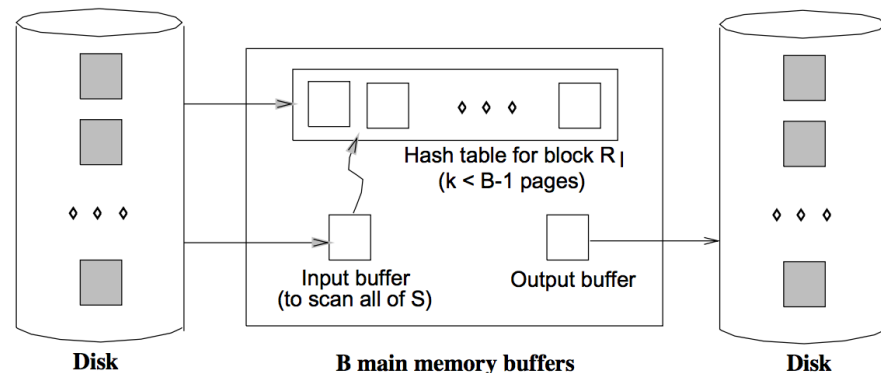
foreach block of $B - 2$ pages of R do

foreach page of S do {

for all matching in-memory tuples $r \in R\text{-block}$ and $s \in S\text{-page}$,
add $\langle r, s \rangle$ to result

}

Relations R and S



Query Processing

Join Operation

2. Block Nested Loops Join

```
SELECT *
FROM   Reserves R, Sailors S
WHERE  R.sid = S.sid
```

```
foreach block of  $B - 2$  pages of  $R$  do
  foreach page of  $S$  do {
    for all matching in-memory tuples  $r \in R\text{-block}$  and  $s \in S\text{-page}$ ,
    add  $\langle r, s \rangle$  to result
  }
```

- cost = M I/Os for reading in R
- S is scanned a total of $\lceil M/B-2 \rceil$ times - ignoring the extra space required per page due to the in-memory hash table, costs = N I/Os
- total cost = $M + N * \lceil M/B-2 \rceil$

assume that we have enough buffers to hold an in-memory hash table for 100 pages of Reserve with at least two additional buffers:

- scan Reserves cost = 1,000 I/Os
- for each 100-page block of Reserves, scan Sailors at 10 scans of Sailors, each costing 500 I/Os
- Total cost = $1,000 + 10 * 500 = 6,000$ I/Os
- if we have buffers to hold only 90 pages of Reserves, scan Sailors $\lceil 1,000/90 \rceil = 12$ times, total cost = $1,000 + 12 * 500 = 7,000$ I/Os

Query Processing

Join Operation

2. Block Nested Loops Join

```
SELECT *
FROM   Reserves R, Sailors S
WHERE  R.sid = S.sid
```

```
foreach block of  $B - 2$  pages of  $R$  do
  foreach page of  $S$  do {
    for all matching in-memory tuples  $r \in R\text{-block}$  and  $s \in S\text{-page}$ ,
    add  $\langle r, s \rangle$  to result
  }
```

- cost = M I/Os for reading in R
- S is scanned a total of $\lceil M/B-2 \rceil$ times - ignoring the extra space required per page due to the in-memory hash table, costs = N I/Os
- total cost = $M + N * \lceil M/B-2 \rceil$

if we choose Sailors to be the outer relation R :

- scan Sailors cost = 500 I/Os
- scan Reserves $\lceil 500/100 \rceil = 5$ times
- total cost is $500 + 5 * 1,000 = 5,500$ I/Os
- if we have buffers to hold only 90 pages of Sailors, scan Reserves $\lceil 500/90 \rceil = 6$ times, total cost = $500 + 6 * 1,000 = 6,500$ I/Os

Query Processing

Join Operation

3. Index Nested Loops Join

```
SELECT *  
FROM   Reserves R, Sailors S  
WHERE  R.sid = S.sid
```

```
foreach tuple  $r \in R$  do  
    foreach tuple  $s \in S$  where  $r_i == s_j$   
        add  $\langle r, s \rangle$  to result
```

- if there is an index on one of the relations on the join attribute(s), make the indexed relation be the inner relation
- index on S: for each tuple $r \in R$, use the index to retrieve matching tuples of S (in the same partition having the same value in the join column) cost depends on:
 1. index on S is a B+ tree index: cost to find leaf is typically 2 to 4 I/Os
index on S is a hash index: cost to find bucket is 1 or 2 I/Os
 2. cost of retrieving matching S tuples depends on whether the index is clustered:
 - clustered: cost per outer tuple $r \in R$ is typically just one more I/O
 - unclustered: cost could be one I/O per matching S-tuple (each can be on a different page)

Query Processing

Join Operation

3. Index Nested Loops Join

```
SELECT *  
FROM   Reserves R, Sailors S  
WHERE  R.sid = S.sid
```

```
foreach tuple  $r \in R$  do  
    foreach tuple  $s \in S$  where  $r_i == s_j$   
        add  $\langle r, s \rangle$  to result
```

- hash-based index on the sid attribute of Sailors takes about 1.2 I/Os on average to retrieve page of the index
- sid is a key for Sailors, so at most one matching tuple
- sid in Reserves is a foreign key - one matching Sailors tuple for each Reserves tuple
- cost of scanning Reserves is 1,000
- $100 * 1,000 = 100,000$ tuples in Reserves
- retrieve the Sailors page containing the qualifying tuple
- total cost = $100,000 * (1 + 1.2) = 221,000$ I/Os

Query Processing

Join Operation

3. Index Nested Loops Join

```
SELECT *  
FROM   Reserves R, Sailors S  
WHERE  R.sid = S.sid
```

```
foreach tuple  $r \in R$  do  
    foreach tuple  $s \in S$  where  $r_i == s_j$   
        add  $\langle r, s \rangle$  to result
```

- hash-based index on the sid attribute of Reserves
- scan Sailors - 500 I/Os
- $80 * 500 = 40,000$ Sailors tuples
- for each tuple, use the index to retrieve matching Reserves tuples (each tuple could match with either zero or more Reserves tuples) in 1.2 I/Os
- total cost = $500 + 40,000 * 1.2 = 48,500$ I/Os + cost of retrieving matching Reserves tuples

Query Processing

Join Operation

3. Index Nested Loops Join

SELECT *	foreach tuple $r \in R$ do
FROM Reserves R, Sailors S	foreach tuple $s \in S$ where $r_i == s_j$
WHERE R.sid = S.sid	add $\langle r, s \rangle$ to result

- total cost = $500 + 40,000 * 1.2 = 48,500$ I/Os + cost of retrieving matching Reserves tuples
- 100,000 reservations for 40,000 Sailors, each Sailors tuple matches with 2.5 Reserves tuples on average
- clustered index on Reserves: matching tuples are on the same page of Reserves, cost = 1 I/O per Sailor tuple = 40,000 extra I/Os and total cost = $48,500 + 40,000 = 88,500$ I/Os
- unclustered index on Reserves: Reserves tuple may be on a different page, cost = $2.5 * 40,000$ I/Os = 100,000 extra I/Os and total cost = $48,500 + 100,000 = 148,500$ I/Os

Query Processing

Join Operation

4. Sort-Merge Join

- sort both relations on the join attribute
- look for qualifying tuples $r \in R$ and $s \in S$ by merging the two relations
- sorting step groups all tuples with the same value in the join column (easy to identify partitions)
- compare R tuples in a partition with only the S tuples in the same partition (rather than with all S tuples)
- scan the relations R and S , look for tuples T_r in R and T_s in S such that $T_{r_i} = T_{s_j}$
- advance the scan of R as long as $R \text{ tuple} < \text{current } S \text{ tuple}$ and advance the scan of S as long as the $S \text{ tuple} < \text{current } R \text{ tuple}$
- for each tuple r in the current R partition, scan all tuples s in the current S partition and output the joined tuple $\langle r, s \rangle$ •

Query Processing

Join Operation

4. Sort-Merge Join

Algorithm:

```

proc smjoin( $R, S, 'R_i = S'_j$ )

  if  $R$  not sorted on attribute  $i$ , sort it;
  if  $S$  not sorted on attribute  $j$ , sort it;

   $Tr$  = first tuple in  $R$ ;                                // ranges over  $R$ 
   $Ts$  = first tuple in  $S$ ;                                // ranges over  $S$ 
   $Gs$  = first tuple in  $S$ ;                                // start of current  $S$ -partition

  while  $Tr \neq eof$  and  $Gs \neq eof$  do {

    while  $Tr_i < Gs_j$  do
       $Tr$  = next tuple in  $R$  after  $Tr$ ;                    // continue scan of  $R$ 

    while  $Tr_i > Gs_j$  do
       $Gs$  = next tuple in  $S$  after  $Gs$                       // continue scan of  $S$ 

     $Ts = Gs$ ;                                              // Needed in case  $Tr_i \neq Gs_j$ 
    while  $Tr_i == Gs_j$  do {                               // process current  $R$  partition
       $Ts = Gs$ ;                                           // reset  $S$  partition scan
      while  $Ts_j == Tr_i$  do {                             // process current  $R$  tuple
        add  $\langle Tr, Ts \rangle$  to result;                  // output joined tuples
         $Ts$  = next tuple in  $S$  after  $Ts$ ;              // advance  $S$  partition scan
      }                                                    // advance scan of  $R$ 
       $Tr$  = next tuple in  $R$  after  $Tr$ ;
    }                                                       // done with current  $R$  partition

     $Gs = Ts$ ;                                              // initialize search for next  $S$  partition

  }

```

Query Processing

Join Operation

4. Sort-Merge Join

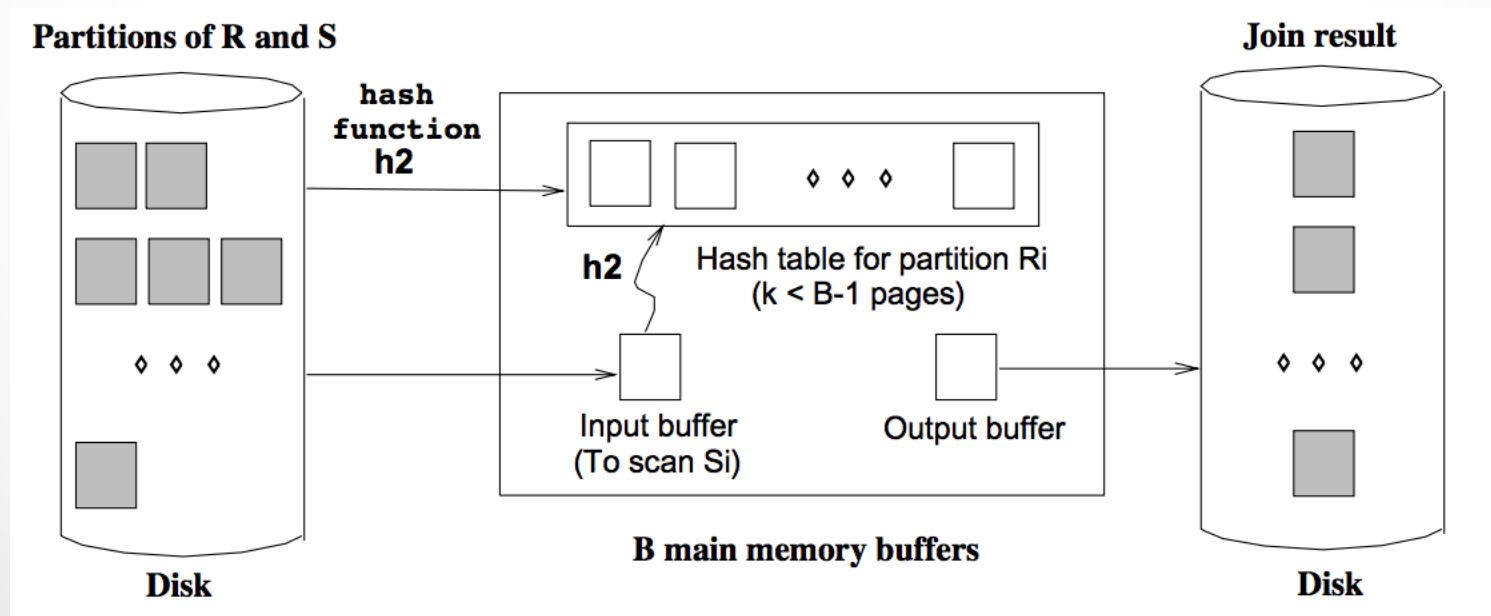
- cost of sorting R is $O(M \log M)$
- cost of sorting S is $O(N \log N)$
- cost of merging phase is $M + N$ (if no S partition is scanned multiple times)
- assume that we have **100** buffer pages
- sort Reserves in 2 passes:
 - first pass produces 10 internally sorted runs of 100 pages each
 - second pass merges these 10 runs to produce the sorted relation
- read and write Reserves in each pass, sorting cost = $2 * 2 * 1000 = 4,000$ I/Os
- sort Sailors in two passes, at a cost of $2 * 2 * 500 = 2,000$ I/Os
- second phase of the sort-merge join algorithm requires an additional scan of both relations
- total cost is $4,000 + 2,000 + 1,000 + 500 = 7,500$ I/Os

Query Processing

Join Operation

5. Hash Join

- **partitioning/building phase:** identify partitions in R and S
- **probing/matching phase:** compare tuples in R partition only with tuples in the corresponding S partition for testing equality join conditions
- hash both relations on the join attribute, using the same hash function h into k partitions
- R tuples in partition i can join only with S tuples in the same partition i



Query Processing

Join Operation

5. Hash Join

Algorithm:

```
// Partition  $R$  into  $k$  partitions
foreach tuple  $r \in R$  do
    read  $r$  and add it to buffer page  $h(r_i)$ ;           // flushed as page fills

// Partition  $S$  into  $k$  partitions
foreach tuple  $s \in S$  do
    read  $s$  and add it to buffer page  $h(s_j)$ ;           // flushed as page fills

// Probing Phase
for  $l = 1, \dots, k$  do {

    // Build in-memory hash table for  $R_l$ , using  $h_2$ 
    foreach tuple  $r \in$  partition  $R_l$  do
        read  $r$  and insert into hash table using  $h_2(r_i)$  ;

    // Scan  $S_l$  and probe for matching  $R_l$  tuples
    foreach tuple  $s \in$  partition  $S_l$  do {
        read  $s$  and probe table using  $h_2(s_j)$ ;
        for matching  $R$  tuples  $r$ , output  $\langle r, s \rangle$  };

    clear hash table to prepare for next partition;
}
```

Query Processing

Join Operation

5. Hash Join

- partitioning phase: scan both R and S once and write them both out once, $\text{cost} = 2(M + N)$
- probing phase: scan each partition once, $\text{cost} = M + N$
- total cost = $3(M + N)$
- assuming each partition fits into memory
- total cost is $3 * (500 + 1,000) = 4,500$ I/Os,

FAU SUMMARY

- no index unsorted file - the only access path is a file scan
- no index but the file is sorted - a binary search to first tuple
- B+ tree index selectivity depends on if the index is clustered or unclustered and the # of result tuples
- hash indexes can be used only for equality selections
- projection operation - implemented by sorting and duplicate elimination during the sorting step
- nested loops join- join condition is evaluated between each pair of tuples from R and S
- block nested loops join performs pairing that minimizes the number of disk accesses
- index nested loops join fetches only matching tuples from S for each tuple of R by using an index
- sort-merge join sorts R and S on the join attributes using an external merge sort and performs pairing during the final merge step
- hash join first partitions R and S using a hash function on the join attributes, only partitions with the same hash values need to be joined in a subsequent step.